

Data-efficient Deep Learning for Earth Observation

Joëlle Hanna, Linus Scheibenreif
University of St. Gallen

Michael Mommert
Stuttgart University of Applied Sciences

What this tutorial is about

What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

In detail, we will focus on two aspects:

What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

In detail, we will focus on two aspects:

- **Data efficiency:** “How can we use available data most efficiently?”

What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

In detail, we will focus on two aspects:

- **Data efficiency:** “How can we use available data most efficiently?”
- **Label efficiency:** “How can we use available labels (or even unlabeled data) most efficiently?”

What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

In detail, we will focus on two aspects:

- **Data efficiency:** “How can we use available data most efficiently?”
- **Label efficiency:** “How can we use available labels (or even unlabeled data) most efficiently?”

We will address these questions in lecture-style presentations of the fundamentals, hands-on coding labs and discussions.

Who we are

Who we are



Joëlle Hanna

PhD student

"Multi-modal Representation
Learning for Remote Sensing"

Who we are



Joëlle Hanna

PhD student
"Multi-modal Representation
Learning for Remote Sensing"



Linus Scheibenreif

PhD student
"Self-supervised Deep Learning
for Earth Observation"

Who we are



Joëlle Hanna

PhD student
"Multi-modal Representation
Learning for Remote Sensing"



Linus Scheibenreif

PhD student
"Self-supervised Deep Learning
for Earth Observation"



University of St.Gallen

Who we are



Joëlle Hanna

PhD student
"Multi-modal Representation
Learning for Remote Sensing"



Linus Scheibenreif

PhD student
"Self-supervised Deep Learning
for Earth Observation"



University of St.Gallen



Michael Mommert

Prof of AI in Remote Sensing
Stuttgart University of Applied Sciences

Hochschule
für Technik
Stuttgart

Today's syllabus

| Time | Content |
|---------------|---|
| 9:00 – 9:15 | Introductions (Michael) |
| 9:15 – 10:00 | Deep Learning Recap and Data Fusion (Michael) |
| 10:00 – 10:15 | Multitask Learning (Joëlle) |
| 10:15 – 10:45 | <i>Coffee Break</i> |
| 10:45 – 11:15 | Multitask Learning (cont'd) (Joëlle) |
| 11:15 – 12:00 | Self-supervised Learning (Linus) |

Resources for this tutorial

- All coding will be done in Jupyter Notebooks. You can access these Notebooks through github: https://github.com/mommermi/IGARSS2024_DataEfficientDeepLearningEO
- We will run the Jupyter Notebooks in the cloud. If possible, we prefer to use Google Colab for this purpose. If you do not have a Google account, please let us know.
- The dataset that we will be using is the ben-ge dataset (see <https://github.com/HSG-AIML/ben-ge> for more information). In this tutorial, we will use a tiny version of ben-ge, which will be made accessible for the time of the tutorial. If you are following this tutorial at some other time, feel free to use the ben-ge-8k dataset (see ben-ge website).

**[https://github.com/mommermi/
IGARSS2024_DataEfficientDeepLearningEO](https://github.com/mommermi/IGARSS2024_DataEfficientDeepLearningEO)**




Data-efficient Deep Learning for Earth Observation

Deep Learning | Data Fusion | Multi-task Learning | Self-Supervised Learning

Deep Learning for Earth observation

Deep Learning for Earth observation


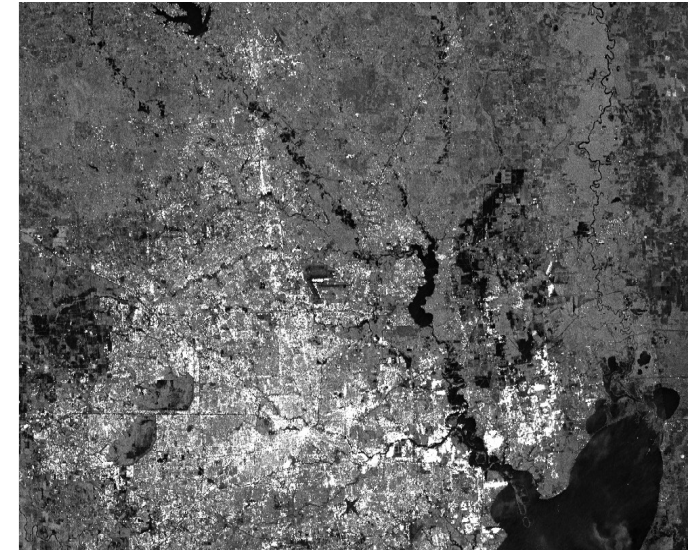
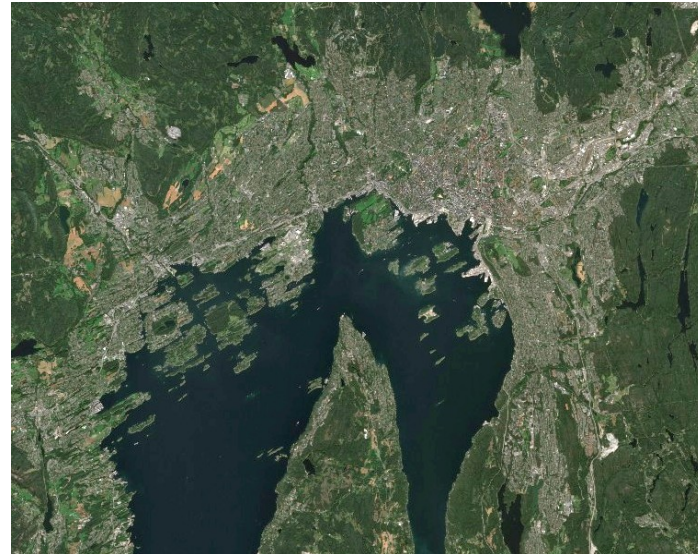
Earth observation data are highly complex
(unstructured, multi-modal).



Deep Learning for Earth observation

Earth observation data are highly complex
(unstructured, multi-modal).

How can we analyze these vast amounts
of data?






Deep Learning for Earth observation

Earth observation data are highly complex
(unstructured, multi-modal).

How can we analyze these vast amounts
of data?

Deep Learning offers the **scalability** to
analyze large amounts of data.



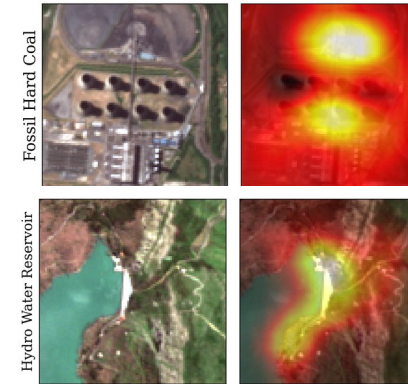
Deep Learning for Earth observation

Earth observation data are highly complex
(unstructured, multi-modal).

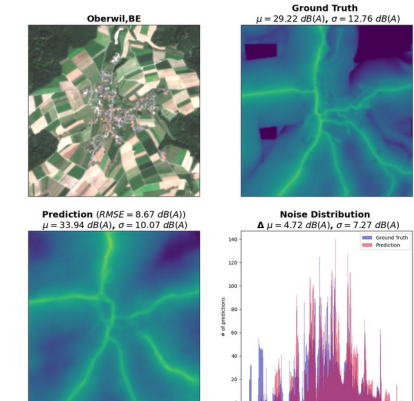
How can we analyze these vast amounts
of data?

Deep Learning offers the **scalability** to
analyze large amounts of data.

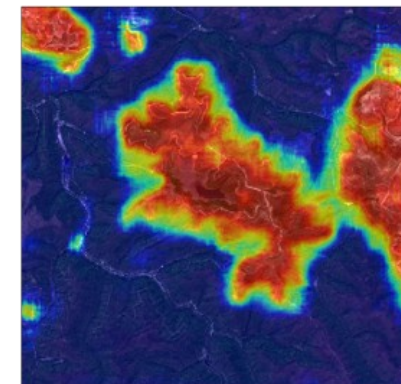
Deep Learning also offers the **flexibility** to
deal with a range of different tasks.



Classification



Regression



Segmentation



**Object
Detection**

Deep Learning for Earth observation

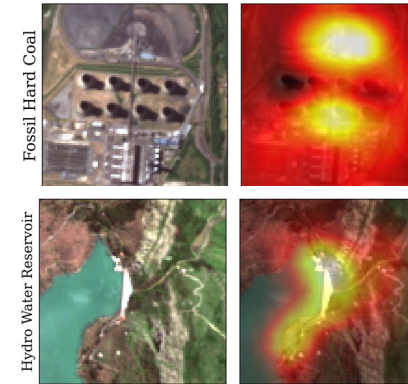
Earth observation data are highly complex
(unstructured, multi-modal).

How can we analyze these vast amounts
of data?

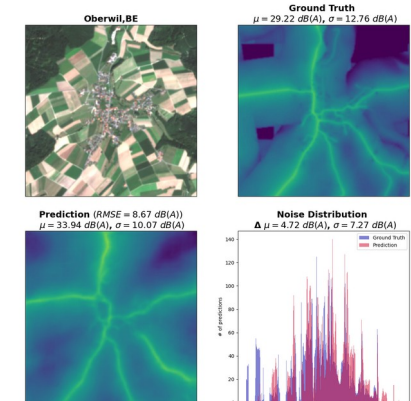
Deep Learning offers the **scalability** to
analyze large amounts of data.

Deep Learning also offers the **flexibility** to
deal with a range of different tasks.


How does it work?



Classification



Regression

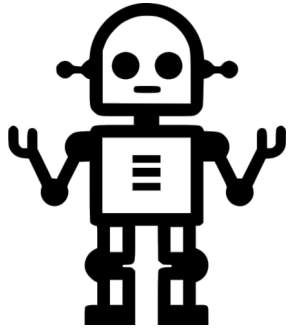


Segmentation




**Object
Detection**

Supervised learning with Neural Networks




Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

Mathematically, it learns a function, f , that maps input data, \mathbf{x} , to the output, \mathbf{y} .


Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

Mathematically, it learns a function, f , that maps input data, x , to the output, y .


Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

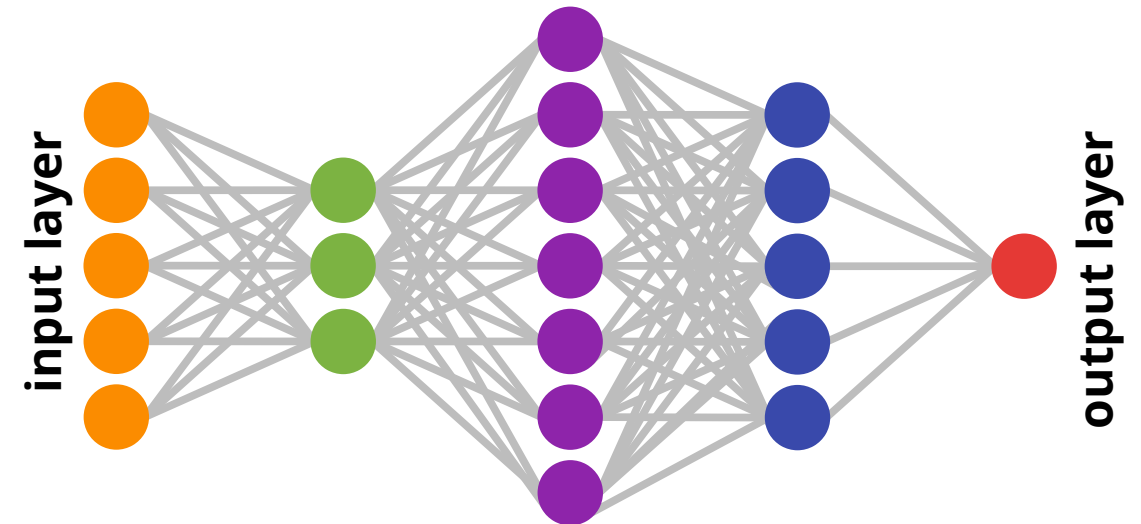
Mathematically, it learns a function, f , that maps input data, x , to the output, y .

Supervised learning with Neural Networks




A machine learns a task from **annotated examples**.

Mathematically, it learns a function, f , that maps input data, x , to the output, y .



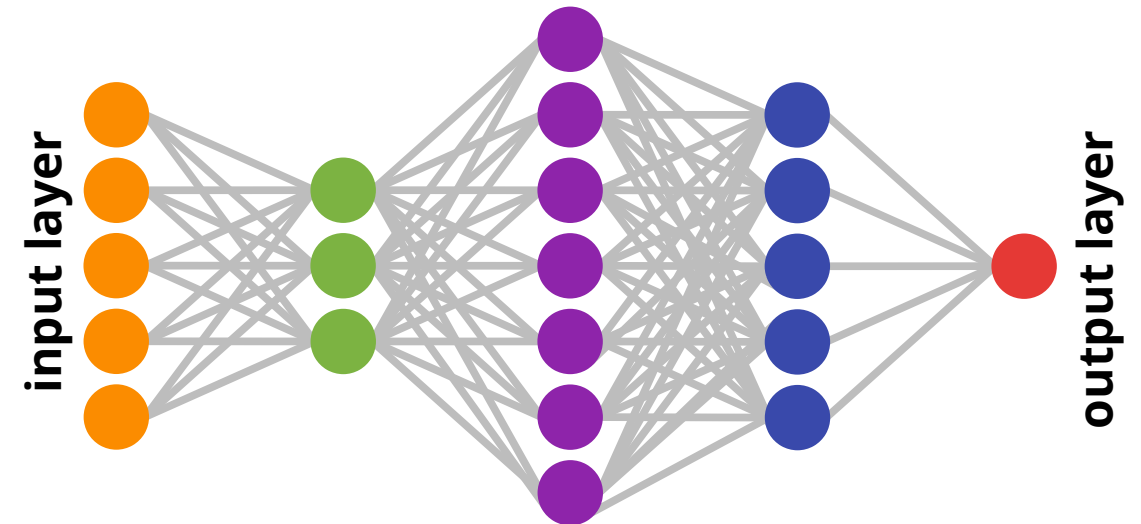
A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.


Mathematically, it learns a function, f , that maps input data, x , to the output, y .



A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.


How does the model learn?

Neural network training pipeline




Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:




Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:
 -) Evaluate model on batch input data (=prediction) in forward pass




Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:
 -) Evaluate model on batch input data (=prediction) in forward pass
 -) Compute loss on prediction and target y




Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:
 -) Evaluate model on batch input data (=prediction) in forward pass
 -) Compute loss on prediction and target y
 -) Compute weight gradients with backprop.




Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:
 -) Evaluate model on batch input data (=prediction) in forward pass
 -) Compute loss on prediction and target y
 -) Compute weight gradients with backprop.
 -) Modify weights based on gradients and learning rate





Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:
 -) Evaluate model on batch input data (=prediction) in forward pass
 -) Compute loss on prediction and target y
 -) Compute weight gradients with backprop.
 -) Modify weights based on gradients and learning rate
 -) Repeat for all batches



Neural network training pipeline



- Sample batch (input data x and target data y) from training dataset:



- Repeat for a number of epochs, monitor training and validation loss + metrics


Neural network training pipeline

- Sample batch (input data x and target data y) from training dataset:




- Repeat for a number of epochs, monitor training and validation loss + metrics
- Stop before overfitting sets in

Supervised learning with Neural Networks




A machine learns a task from **annotated examples**.

Mathematically, it learns a function, f , that maps input data, x , to the output, y .



A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.


Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

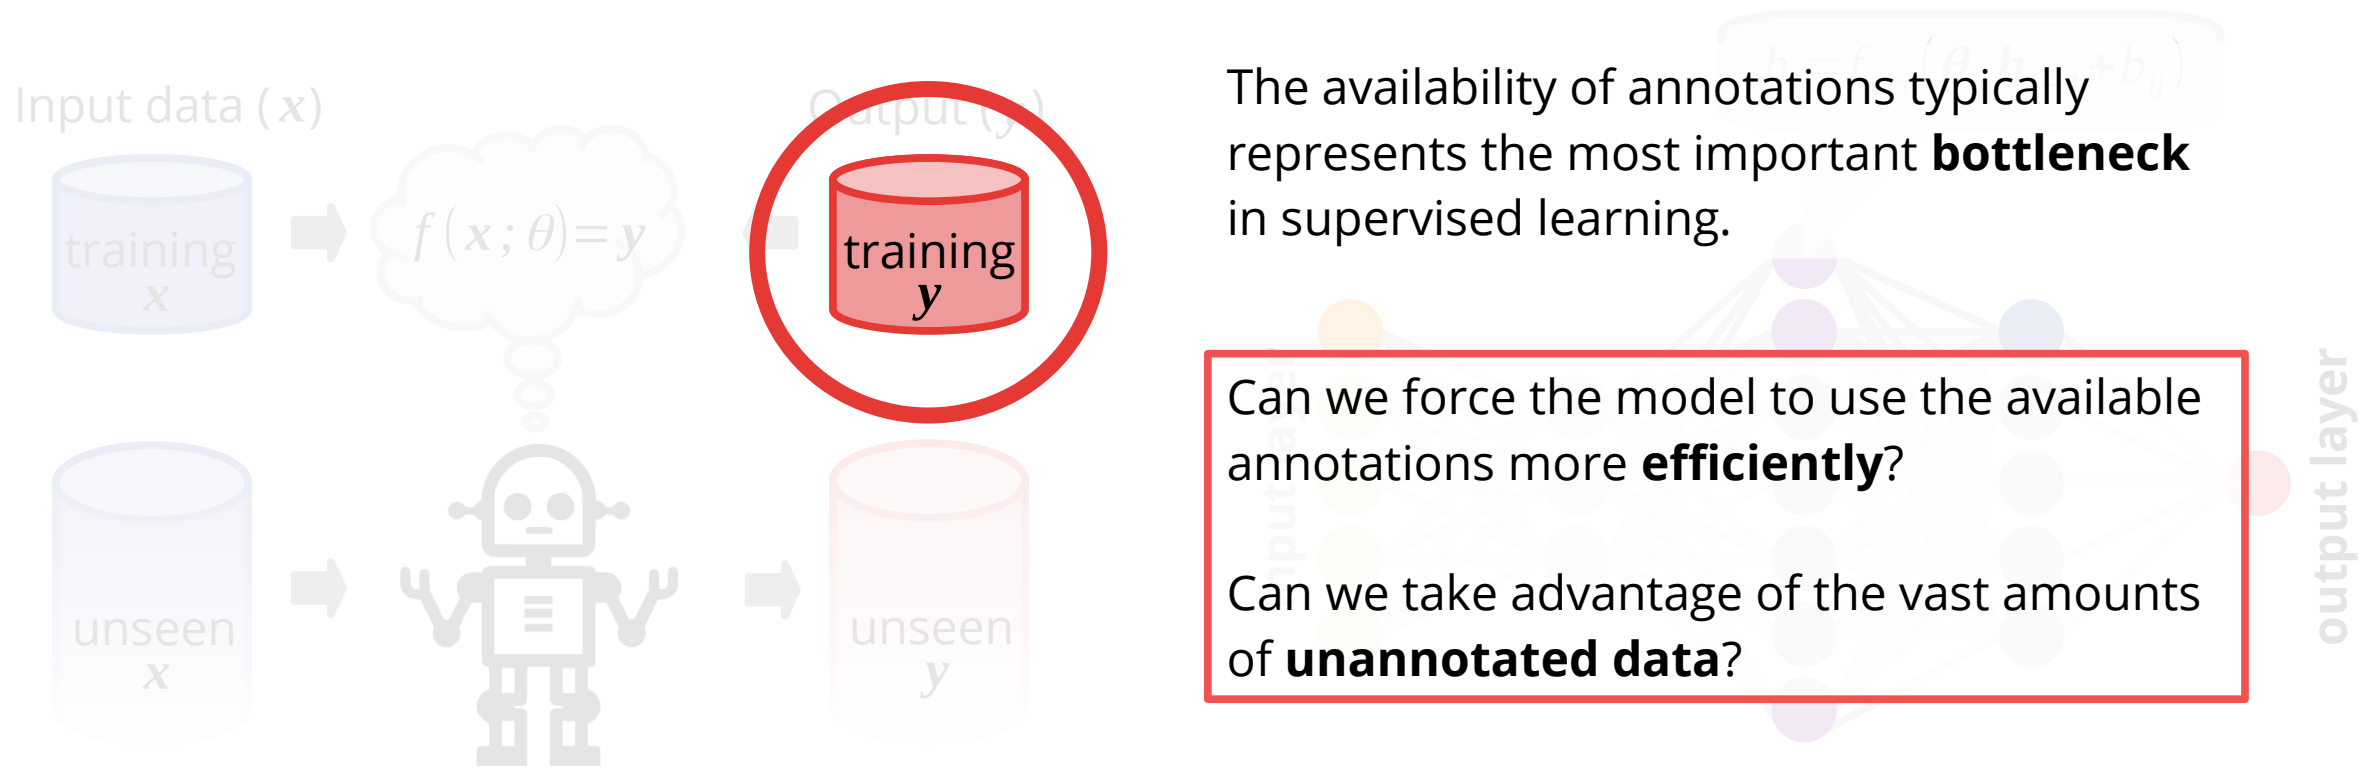
Mathematically, it learns a function, f , that maps input data, x , to the output, y .

The availability of annotations typically represents the most important **bottleneck** in supervised learning.



A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.


Mathematically, it learns a function, f , that maps input data, x , to the output, y .

A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

How can we use annotated data more efficiently?

Before we answer this question, let's have a look at how to implement the

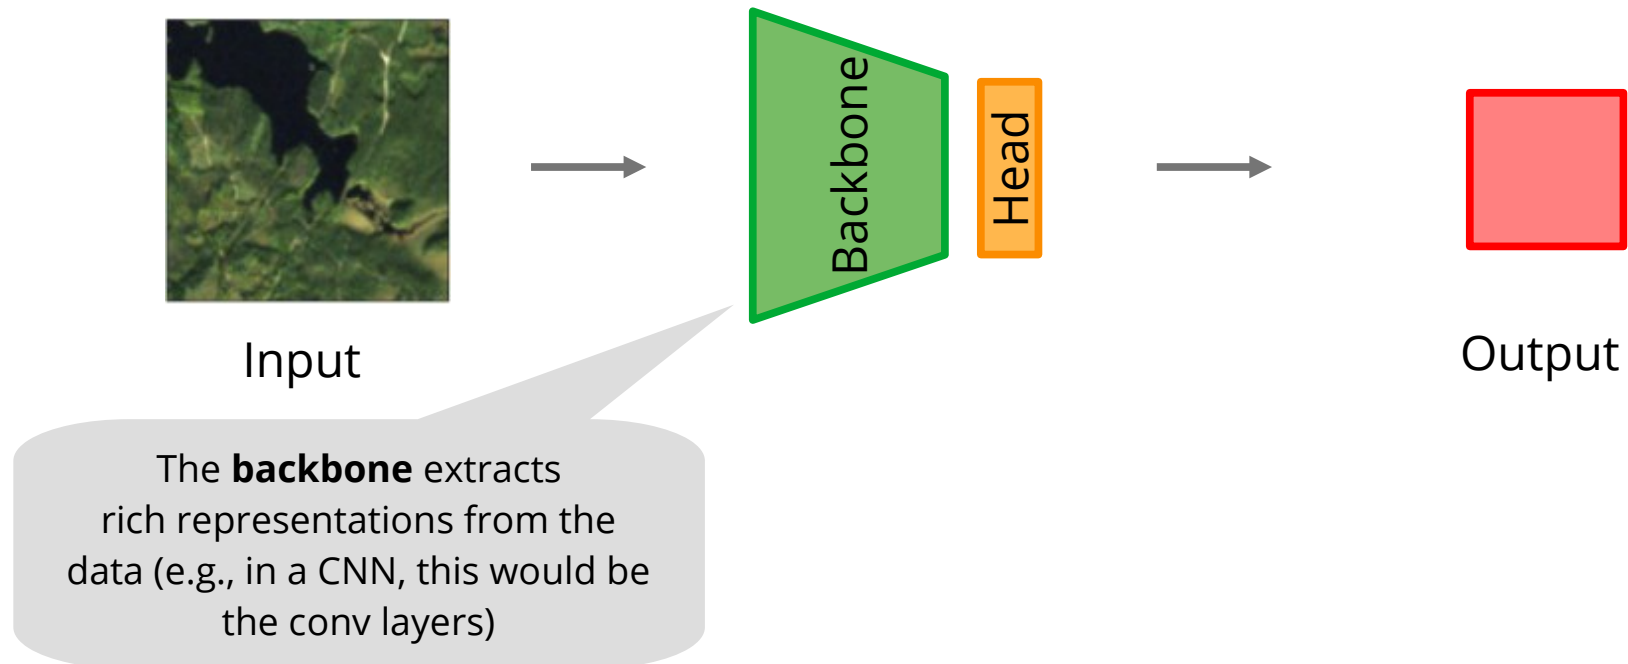
Supervised Learning Setup



How can we use annotated data more efficiently?

Before we answer this question, let's have a look at how to implement the

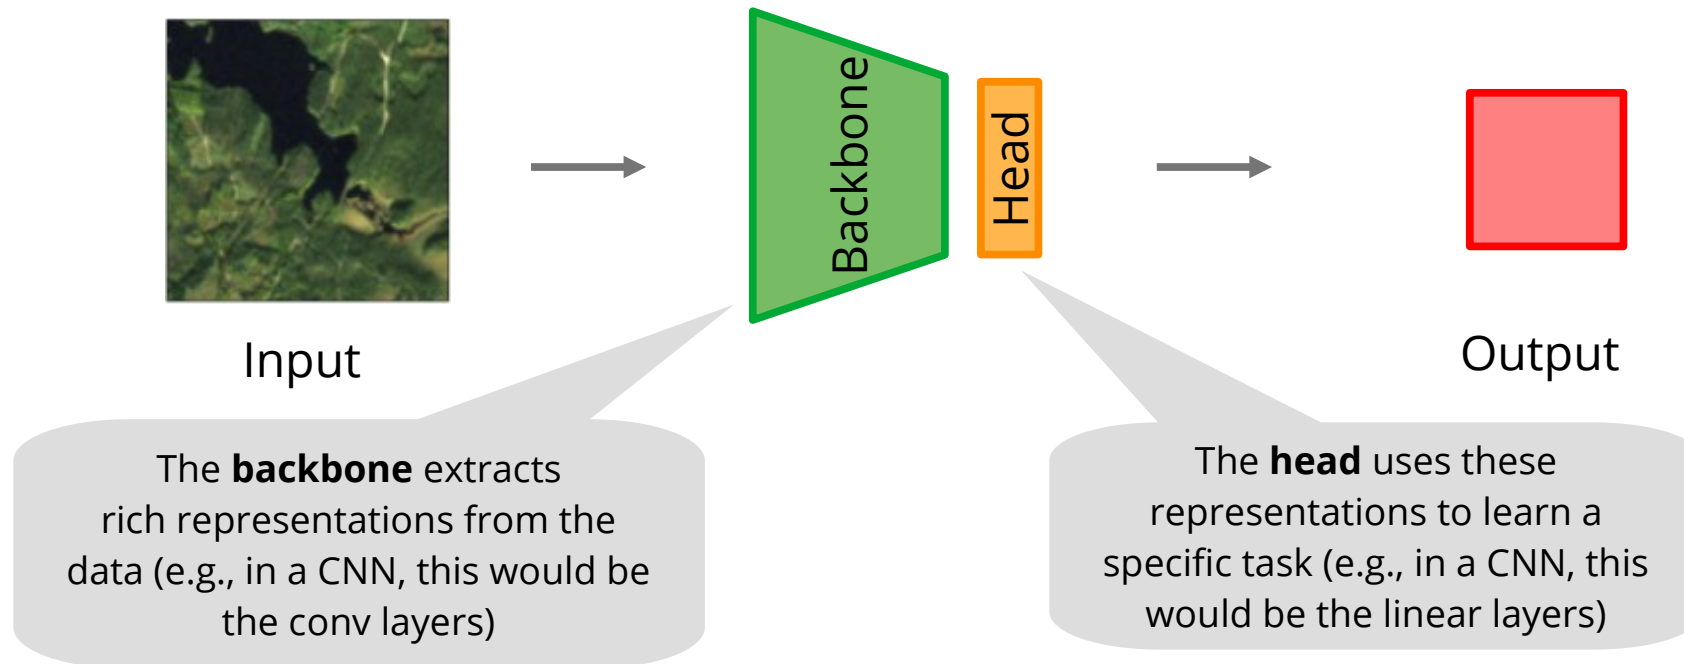
Supervised Learning Setup



How can we use annotated data more efficiently?

Before we answer this question, let's have a look at how to implement the

Supervised Learning Setup



How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.

How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.

Early Data Fusion

How can we use annotated data more efficiently?


We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.

Early Data Fusion

e.g., multi-
spectral data

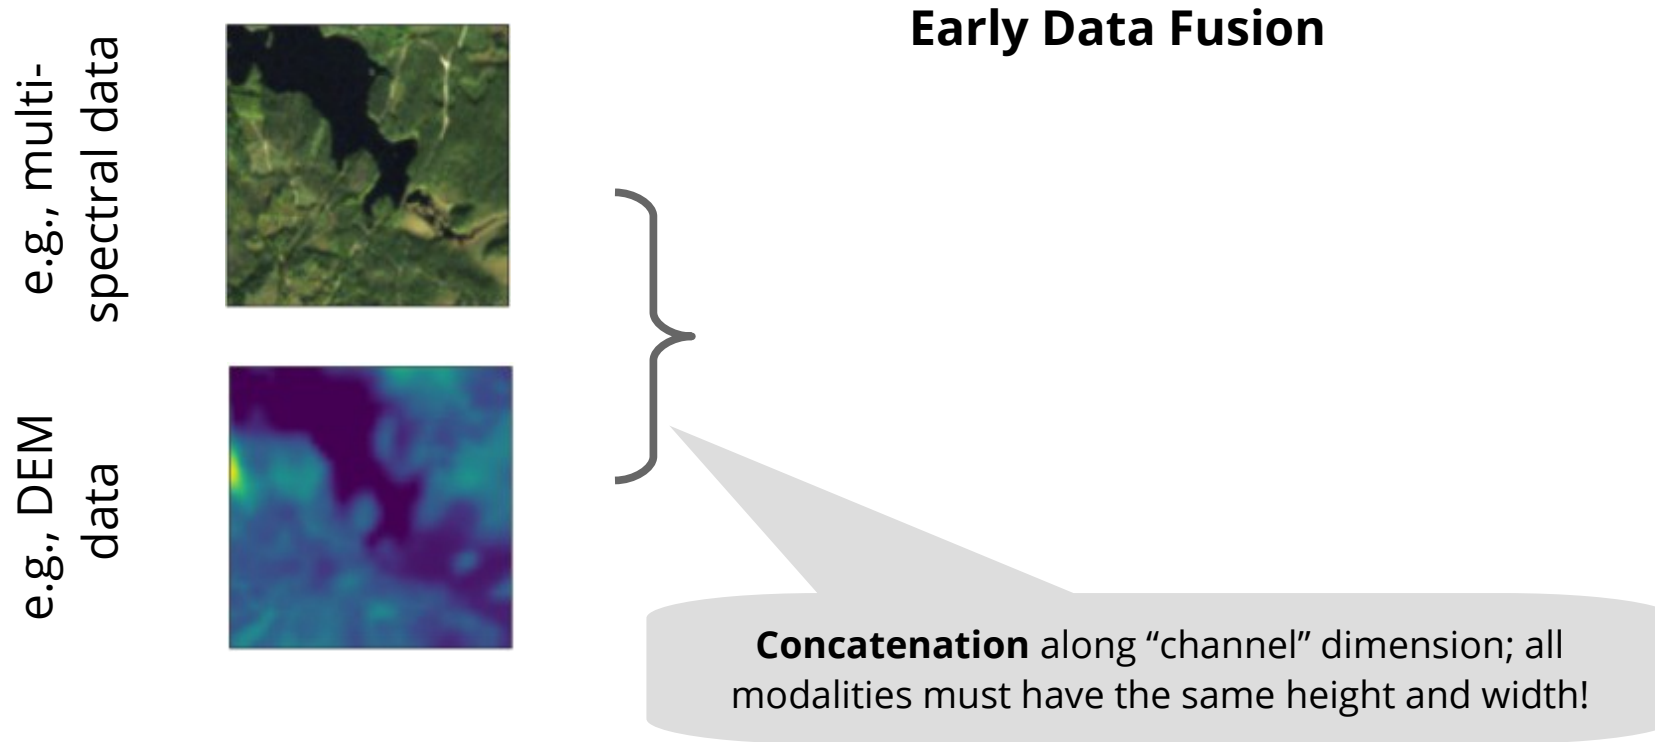


e.g., DEM
data



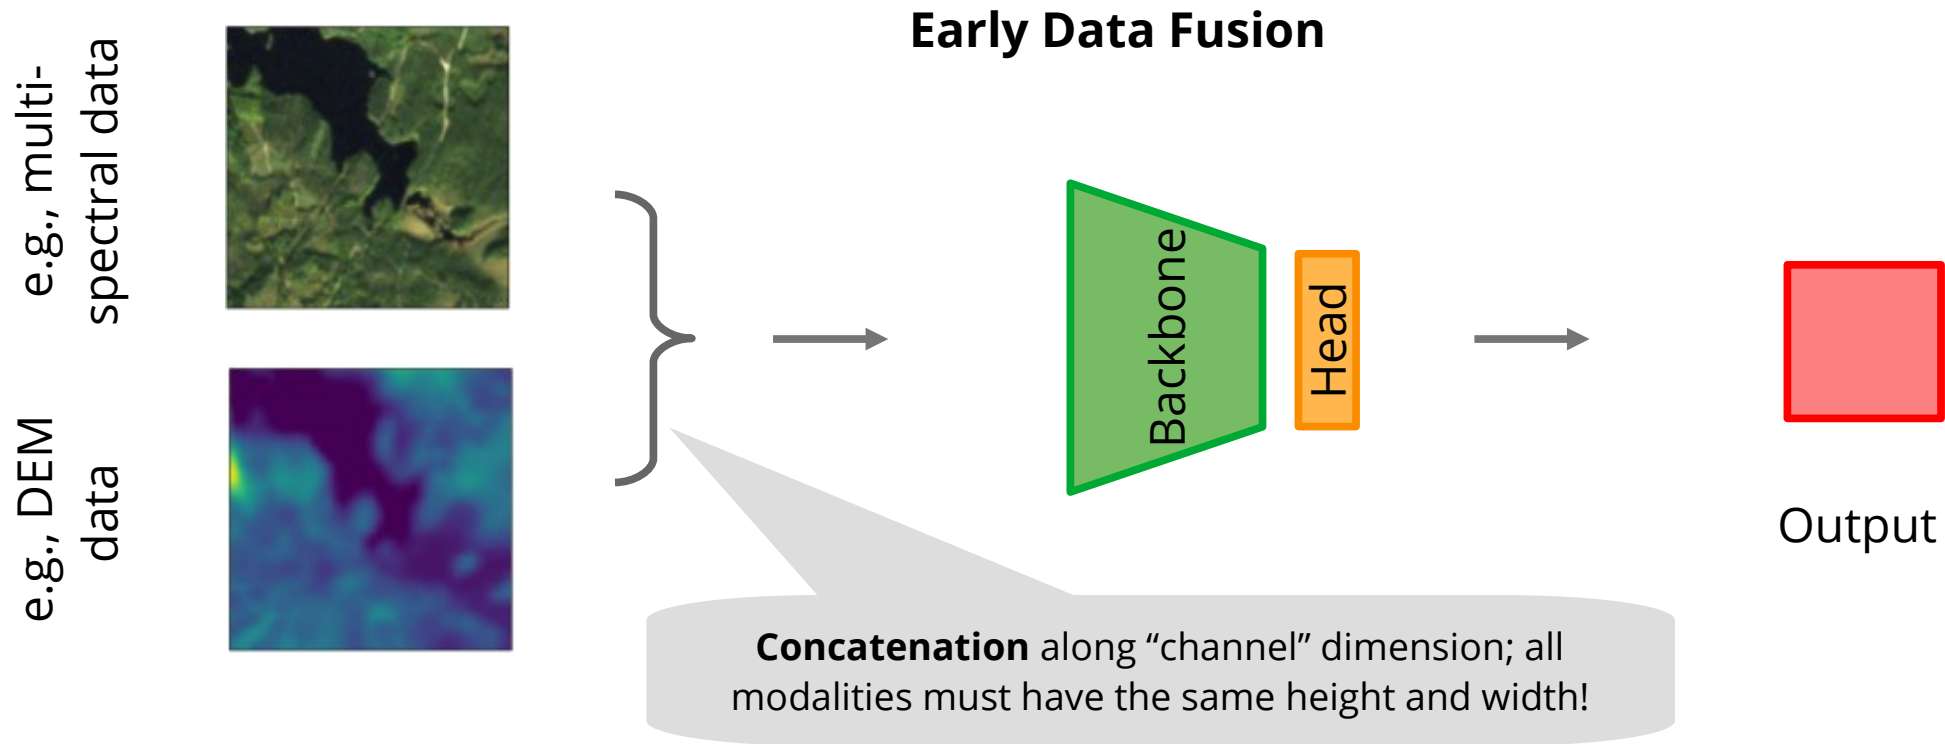
How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.




How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.



How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.



By combining different data modalities early on, we provide additional information to our model.


How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.

e.g., multi-
spectral data




e.g., DEM
data



Late Data Fusion

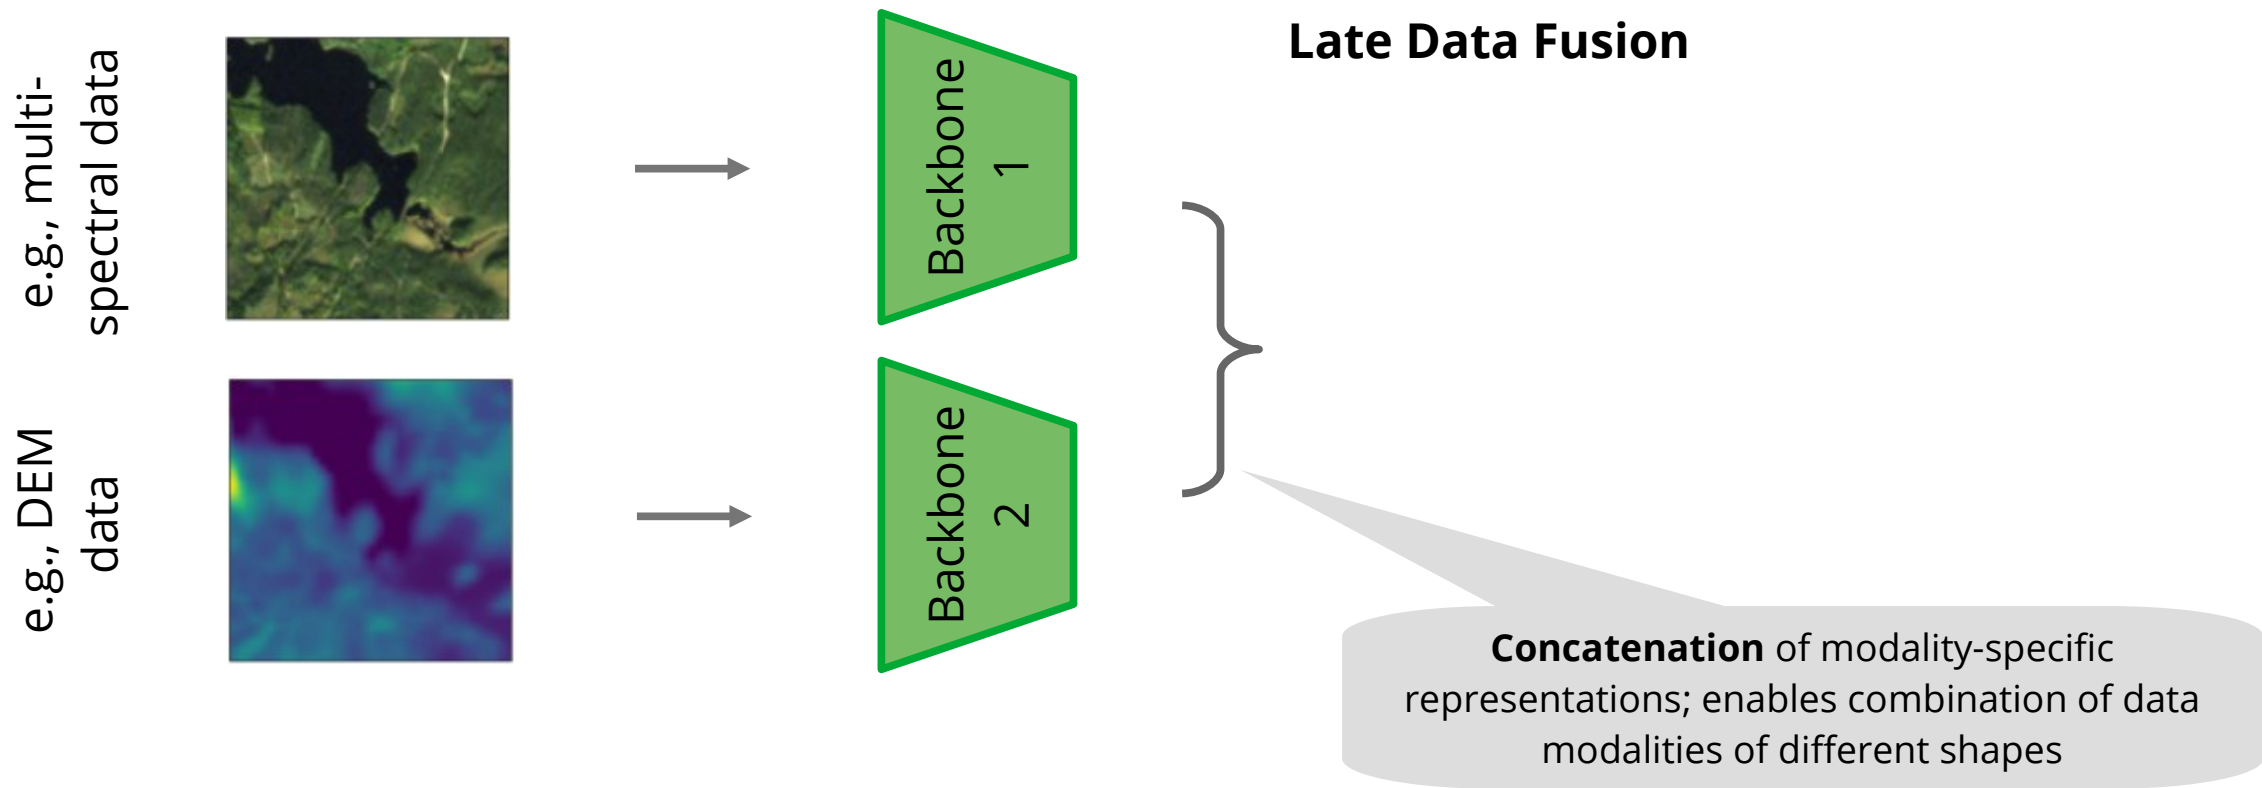
How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.




How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.



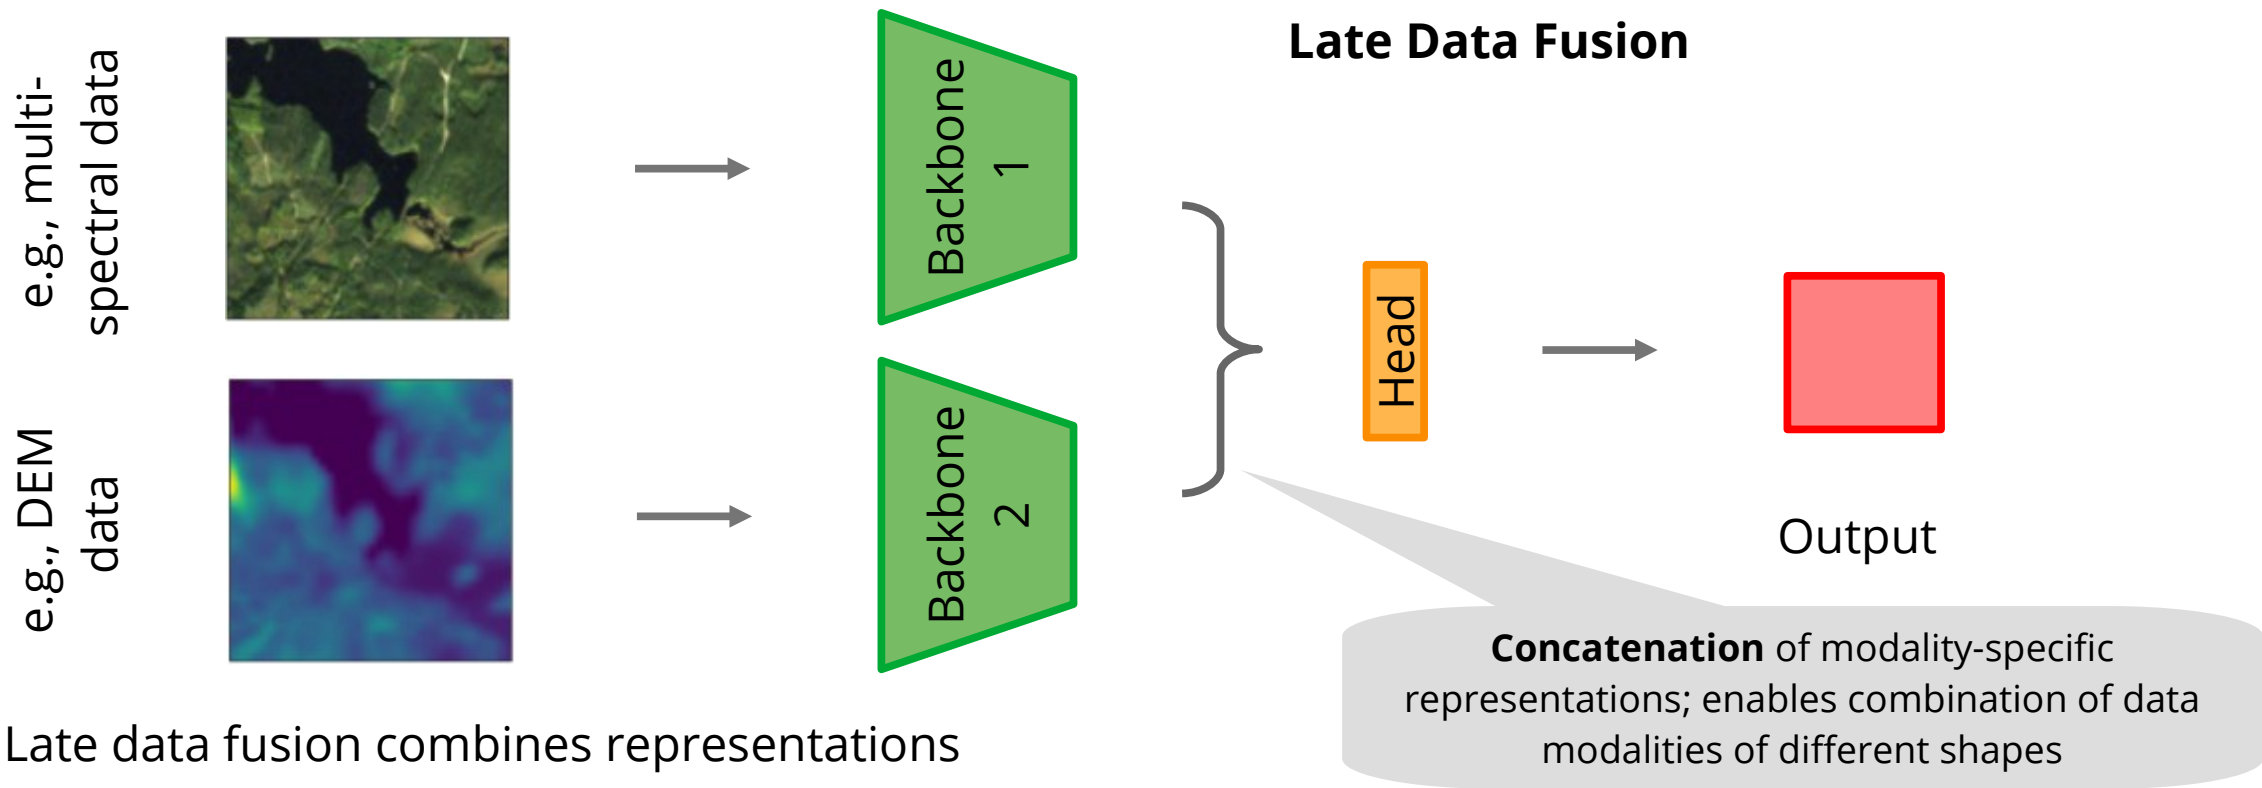
How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.



How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.




Late data fusion combines representations from different data modalities, representing a more flexible approach for data fusion.

How can we use annotated data more efficiently?

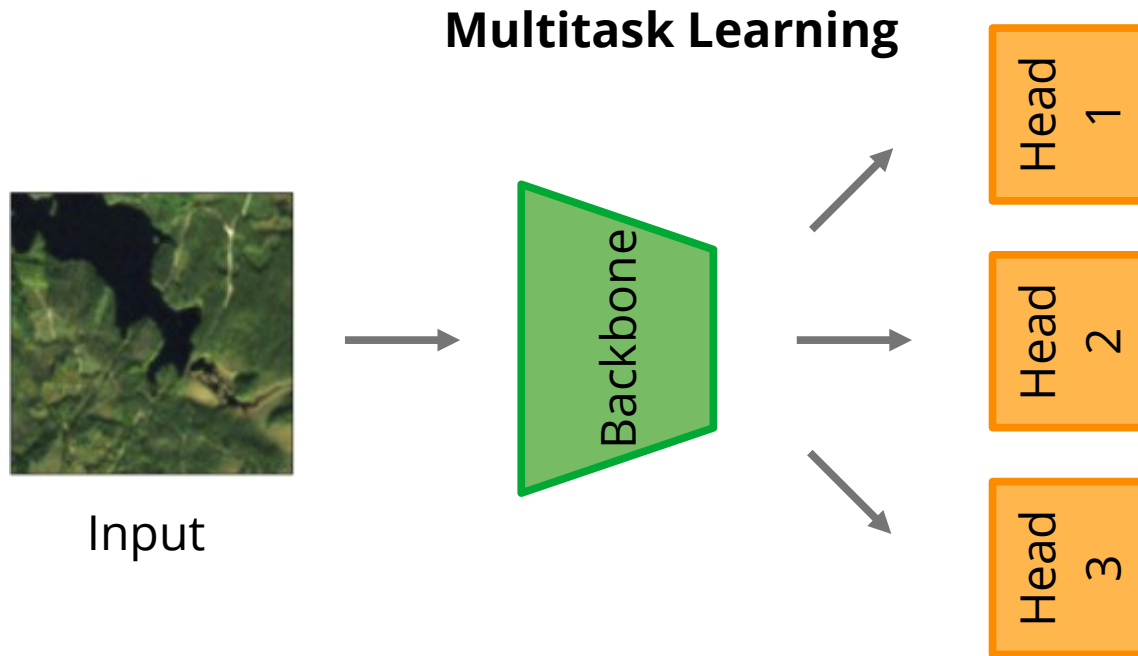
By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.

Multitask Learning




How can we use annotated data more efficiently?

By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.



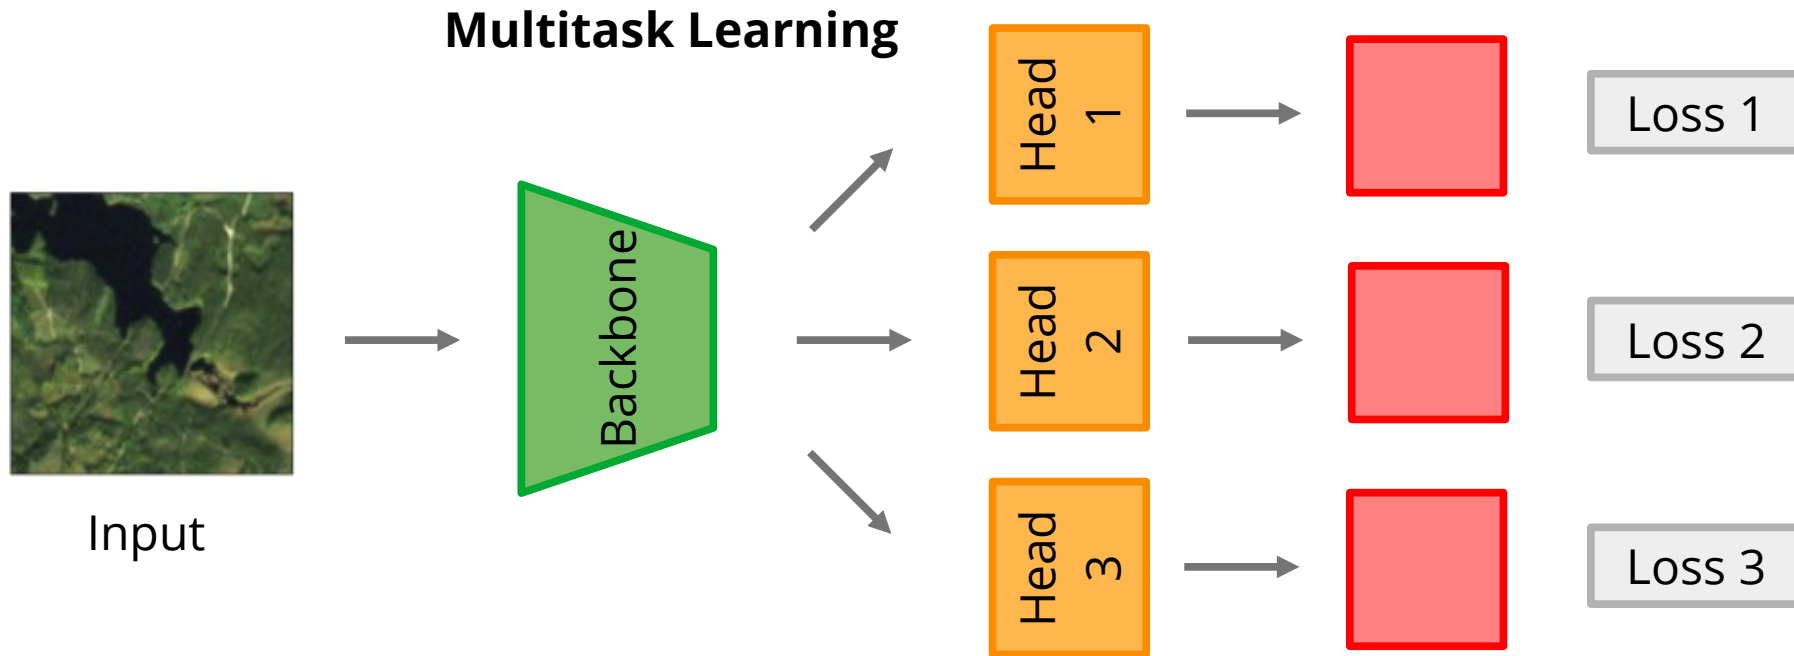
How can we use annotated data more efficiently?

By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.




How can we use annotated data more efficiently?

By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.




How can we use annotated data more efficiently?

By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.



How can we use annotated data more efficiently?


By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.



The different heads are trained simultaneously by optimizing the (weighted) sum of the individual losses. As a result, the performance on each task is (typically) better than if trained individually.


How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



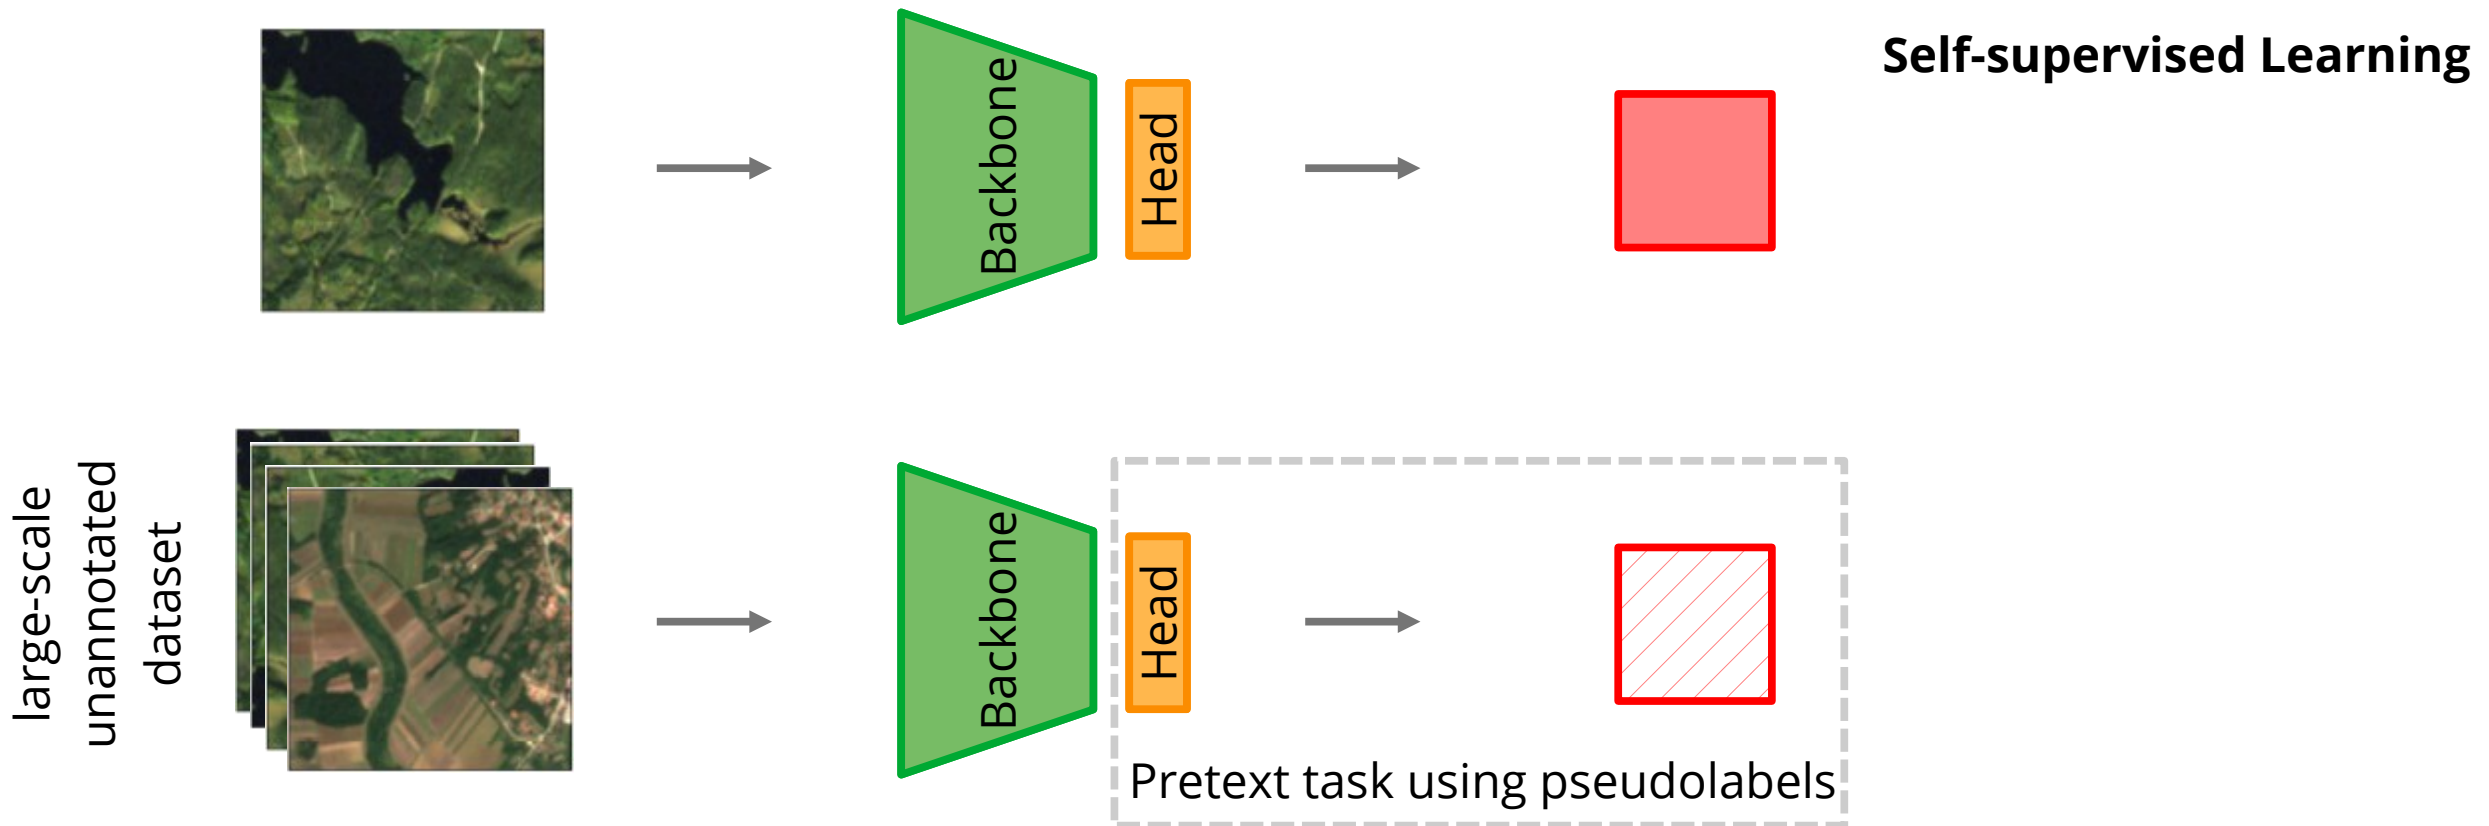
How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



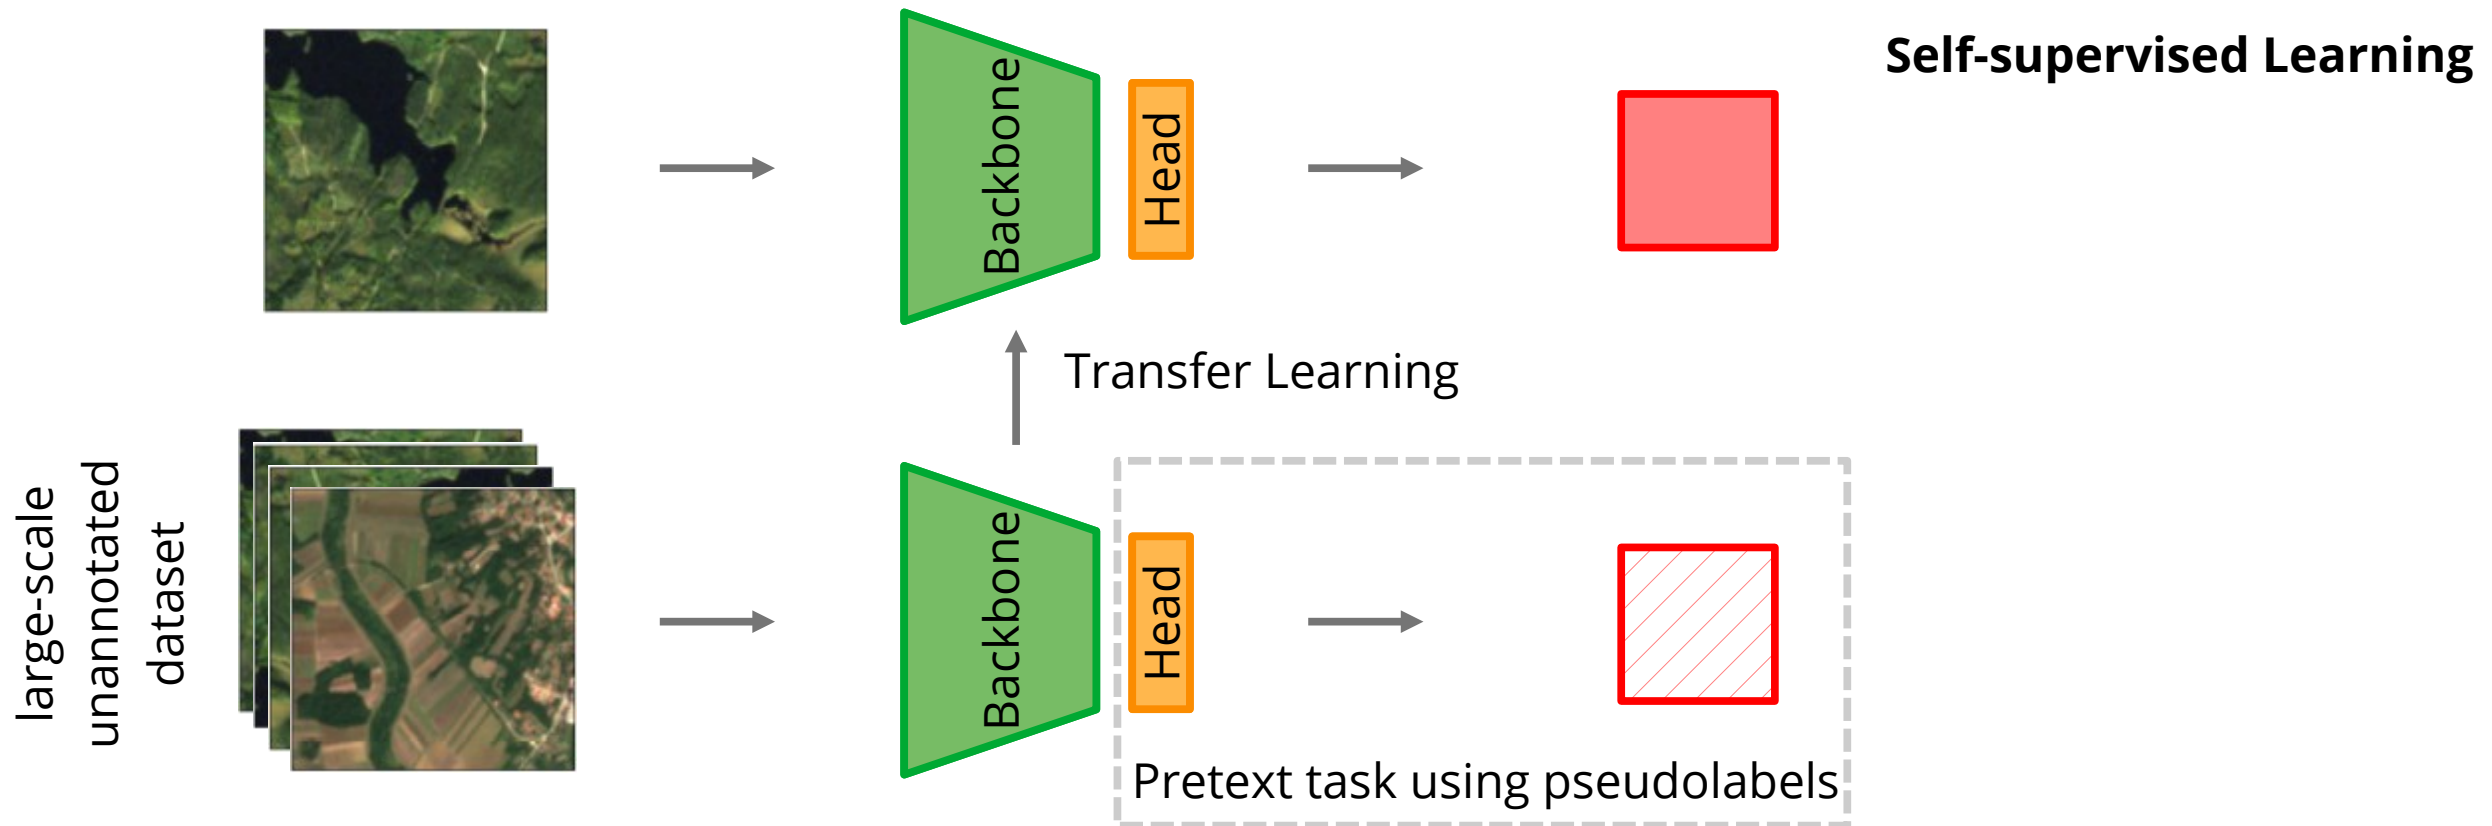
How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



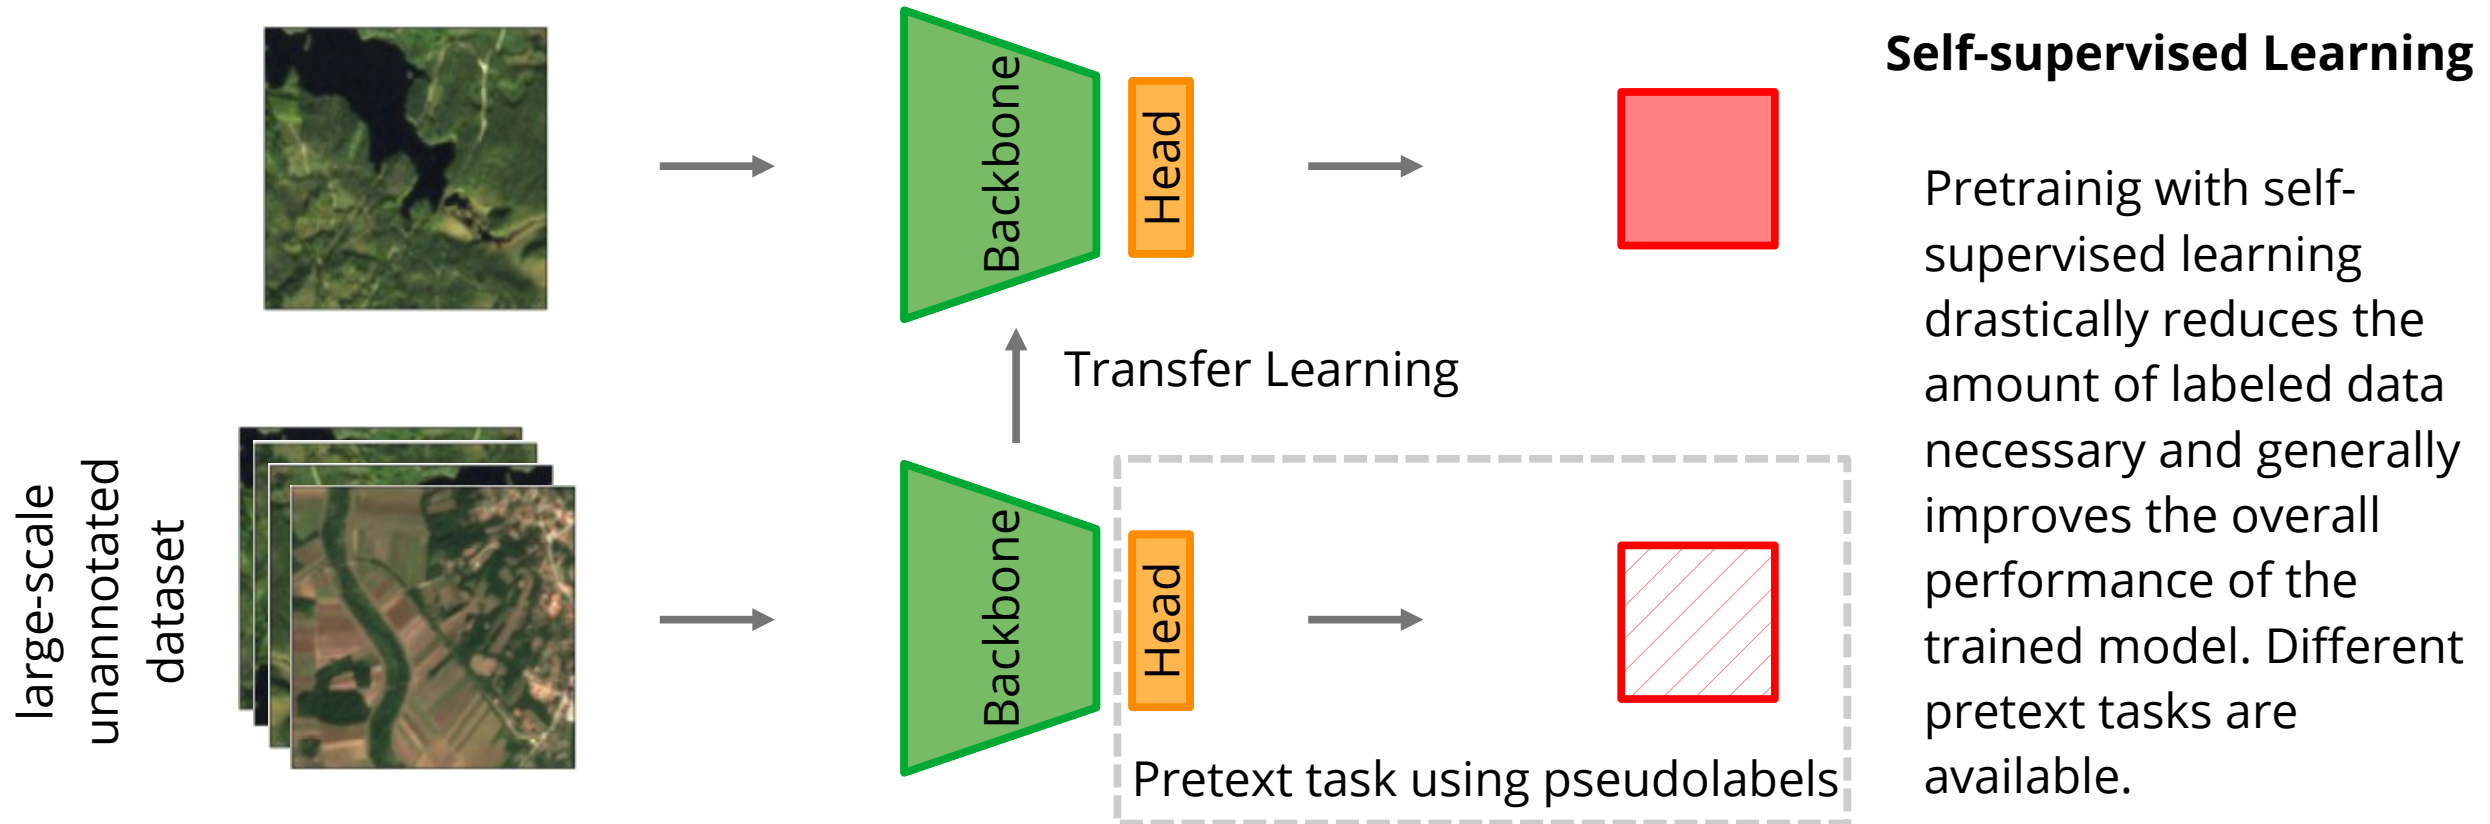
How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



Let's get our hands dirty

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)
- Multitask Learning (Joëlle)

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)
- Multitask Learning (Joëlle)
- Self-supervised Learning (Linus)

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)
- Multitask Learning (Joëlle)
- Self-supervised Learning (Linus)

We will discuss these techniques in separate Jupyter Notebooks. The overall structure and methods for data handling and other aspects should be consistent between these Notebooks for easier understanding.

Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)
- Multitask Learning (Joëlle)
- Self-supervised Learning (Linus)

We will discuss these techniques in separate Jupyter Notebooks. The overall structure and methods for data handling and other aspects should be consistent between these Notebooks for easier understanding.

Feel free to use the code from these Notebooks for your own research!