

# Data-efficient Deep Learning for Earth Observation



Joëlle Hanna, Linus Scheibenreif  
University of St. Gallen



Universität St.Gallen

Michael Mommert  
Stuttgart University of Applied Sciences

Hochschule  
für Technik  
Stuttgart

# What this tutorial is about

As part of this tutorial, we will introduce and discuss different techniques to make more efficient use of data in Deep Learning for Earth observation.

In detail, we will focus on two aspects:

- **Data efficiency:** “How can we use available data most efficiently?”
- **Label efficiency:** “How can we use available labels (or even unlabeled data) most efficiently?”

We will address these questions in lecture-style presentations of the fundamentals, hands-on coding labs and discussions.

# Who we are



**Joëlle Hanna**

PhD student  
"Multi-modal Representation  
Learning for Remote Sensing"



**Linus Scheibenreif**

PhD student  
"Self-supervised Deep Learning  
for Earth Observation"



University of St.Gallen



**Michael Mommert**

Prof of AI in Remote Sensing  
Stuttgart University of Applied Sciences

Hochschule  
für Technik  
Stuttgart

# Today's syllabus

Time	Content
9:00 – 9:15	<b>Introductions</b> (Michael)
9:15 – 10:00	<b>Deep Learning Recap and Data Fusion</b> (Michael)
10:00 – 10:15	<b>Multitask Learning</b> (Joëlle)
10:15 – 10:45	<i>Coffee Break</i>
10:45 – 11:15	<b>Multitask Learning</b> (cont'd) (Joëlle)
11:15 – 12:00	<b>Self-supervised Learning</b> (Linus)

## Resources for this tutorial

- All coding will be done in Jupyter Notebooks. You can access these Notebooks through github: [https://github.com/mommermi/IGARSS2024\\_DataEfficientDeepLearningEO](https://github.com/mommermi/IGARSS2024_DataEfficientDeepLearningEO)
- We will run the Jupyter Notebooks in the cloud. If possible, we prefer to use Google Colab for this purpose. If you do not have a Google account, please let us know.
- The dataset that we will be using is the ben-ge dataset (see <https://github.com/HSG-AIML/ben-ge> for more information). In this tutorial, we will use a tiny version of ben-ge, which will be made accessible for the time of the tutorial. If you are following this tutorial at some other time, feel free to use the ben-ge-8k dataset (see ben-ge website).

**[https://github.com/mommermi/IGARSS2024\\_DataEfficientDeepLearningEO](https://github.com/mommermi/IGARSS2024_DataEfficientDeepLearningEO)**

# **Data-efficient Deep Learning for Earth Observation**




**Deep Learning | Data Fusion | Multi-task Learning | Self-Supervised Learning**

# Deep Learning for Earth observation

Earth observation data are highly complex  
(unstructured, multi-modal).

How can we analyze these vast amounts  
of data?

Deep Learning offers the **scalability** to  
analyze large amounts of data.





# Deep Learning for Earth observation

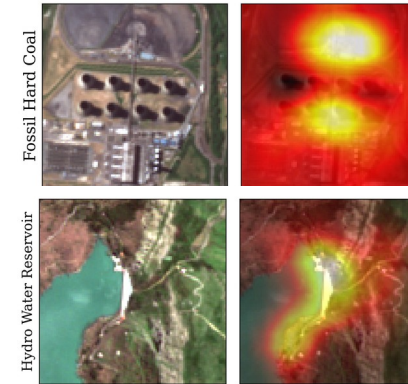
Earth observation data are highly complex  
(unstructured, multi-modal).

How can we analyze these vast amounts  
of data?


Deep Learning offers the **scalability** to  
analyze large amounts of data.

Deep Learning also offers the **flexibility** to  
deal with a range of different tasks.


How does it work?



**Classification**



**Regression**




**Segmentation**



**Object  
Detection**

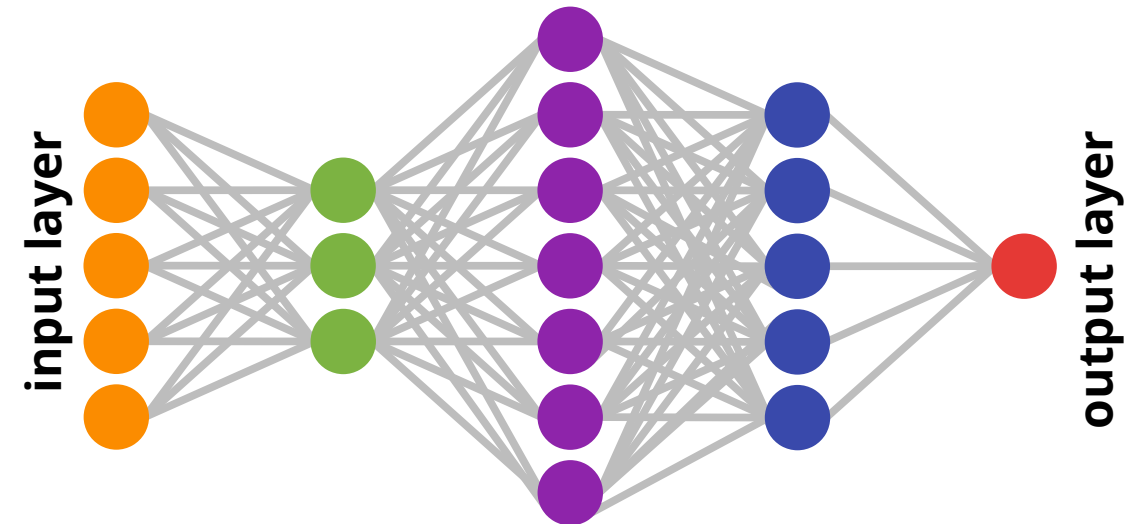


# Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

Mathematically, it learns a function,  $f$ , that maps input data,  $x$ , to the output,  $y$ .


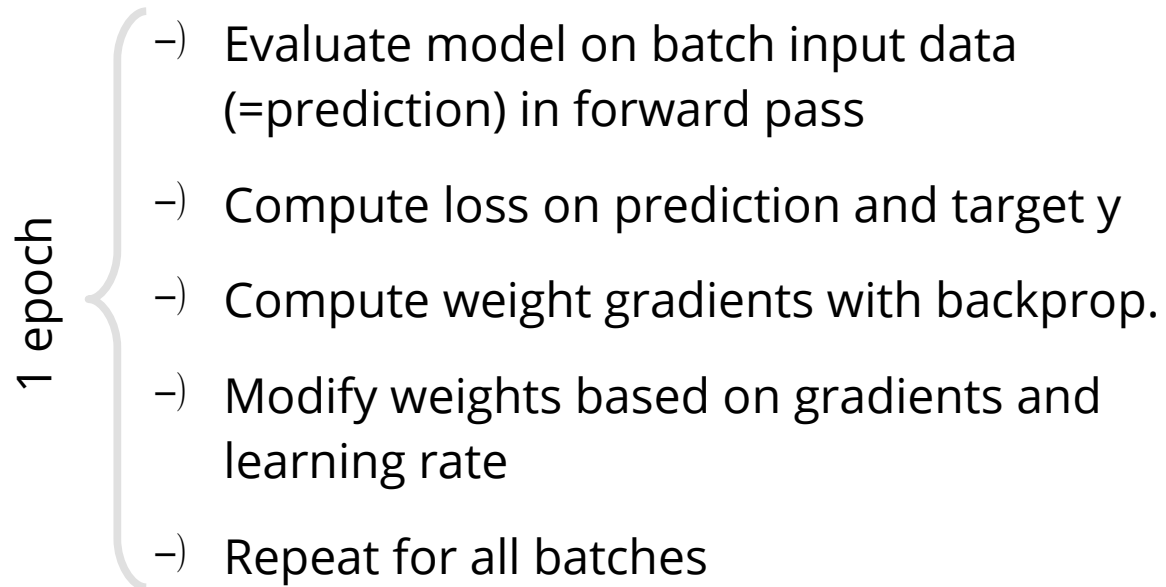


A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

**How does the model learn?**


# Neural network training pipeline

- Sample batch (input data  $x$  and target data  $y$ ) from training dataset:



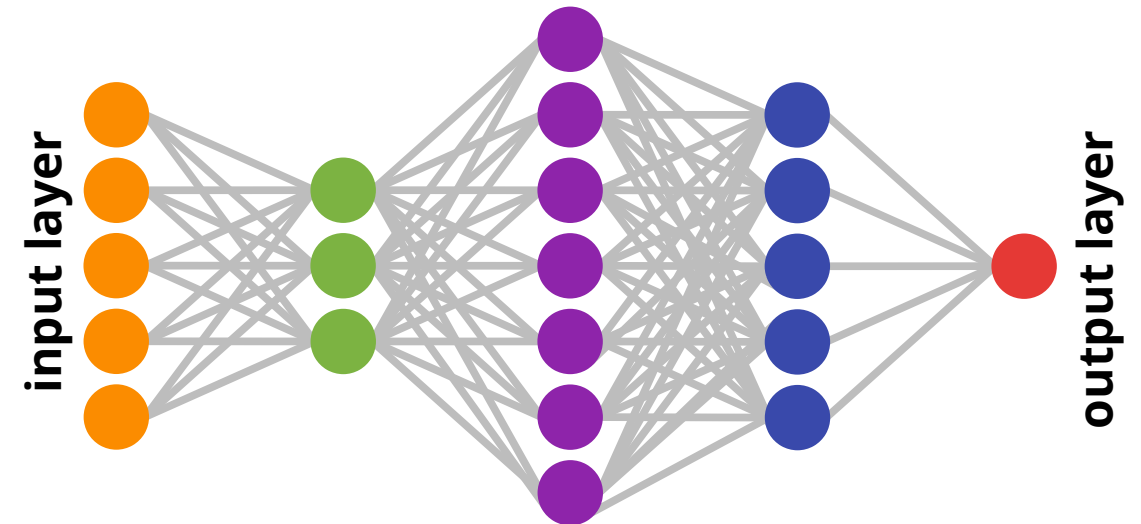
- Repeat for a number of epochs, monitor training and validation loss + metrics
- Stop before overfitting sets in

# Supervised learning with Neural Networks



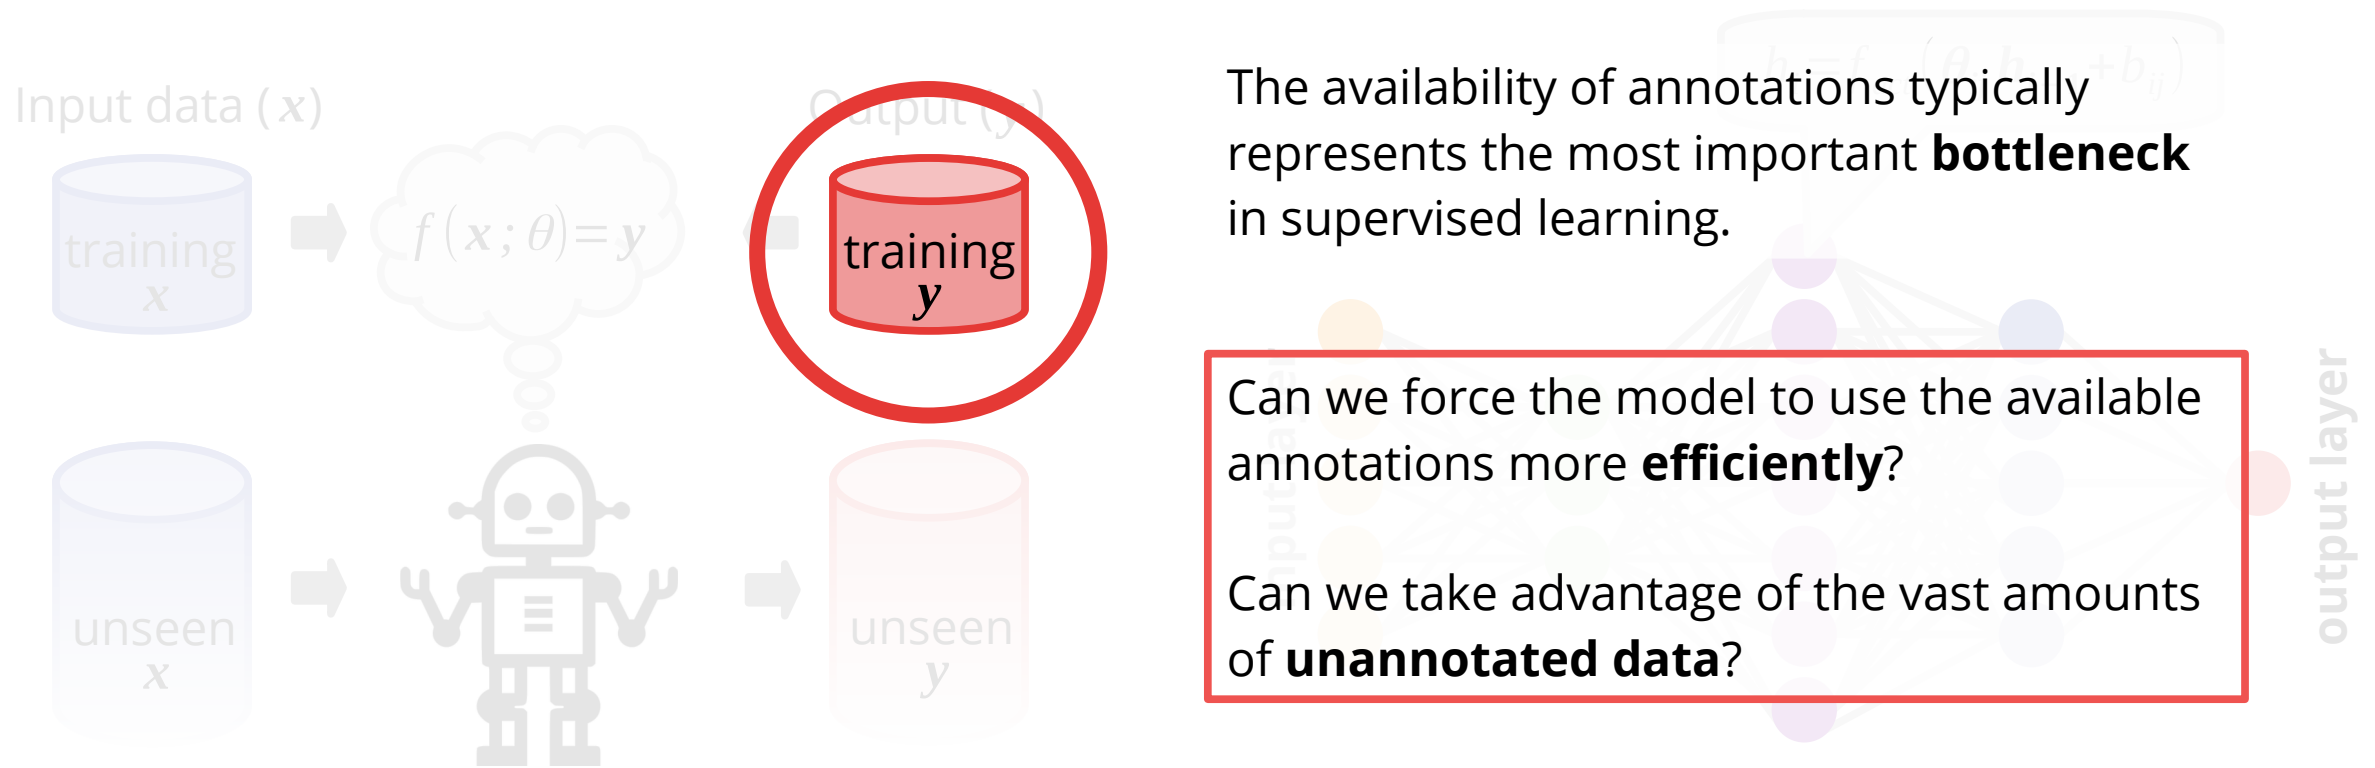
A machine learns a task from **annotated examples**.

Mathematically, it learns a function,  $f$ , that maps input data,  $x$ , to the output,  $y$ .



A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

# Supervised learning with Neural Networks



A machine learns a task from **annotated examples**.

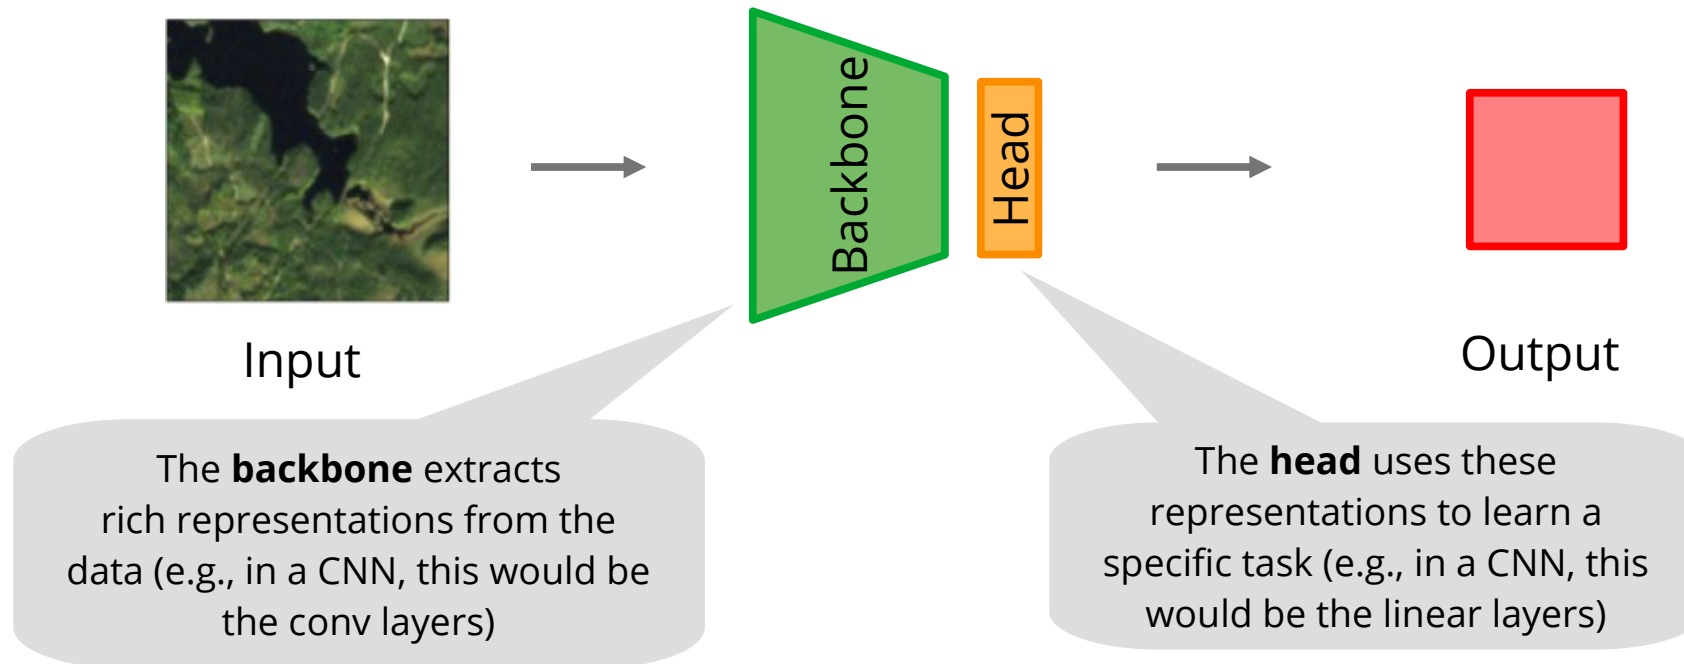
Mathematically, it learns a function,  $f$ , that maps input data,  $x$ , to the output,  $y$ .

A Neural Network is a cascade of mathematical functions; each neuron contains learnable weights that represent the learned knowledge.

# How can we use annotated data more efficiently?


Before we answer this question, let's have a look at how to implement the

## Supervised Learning Setup



# How can we use annotated data more efficiently?

We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.

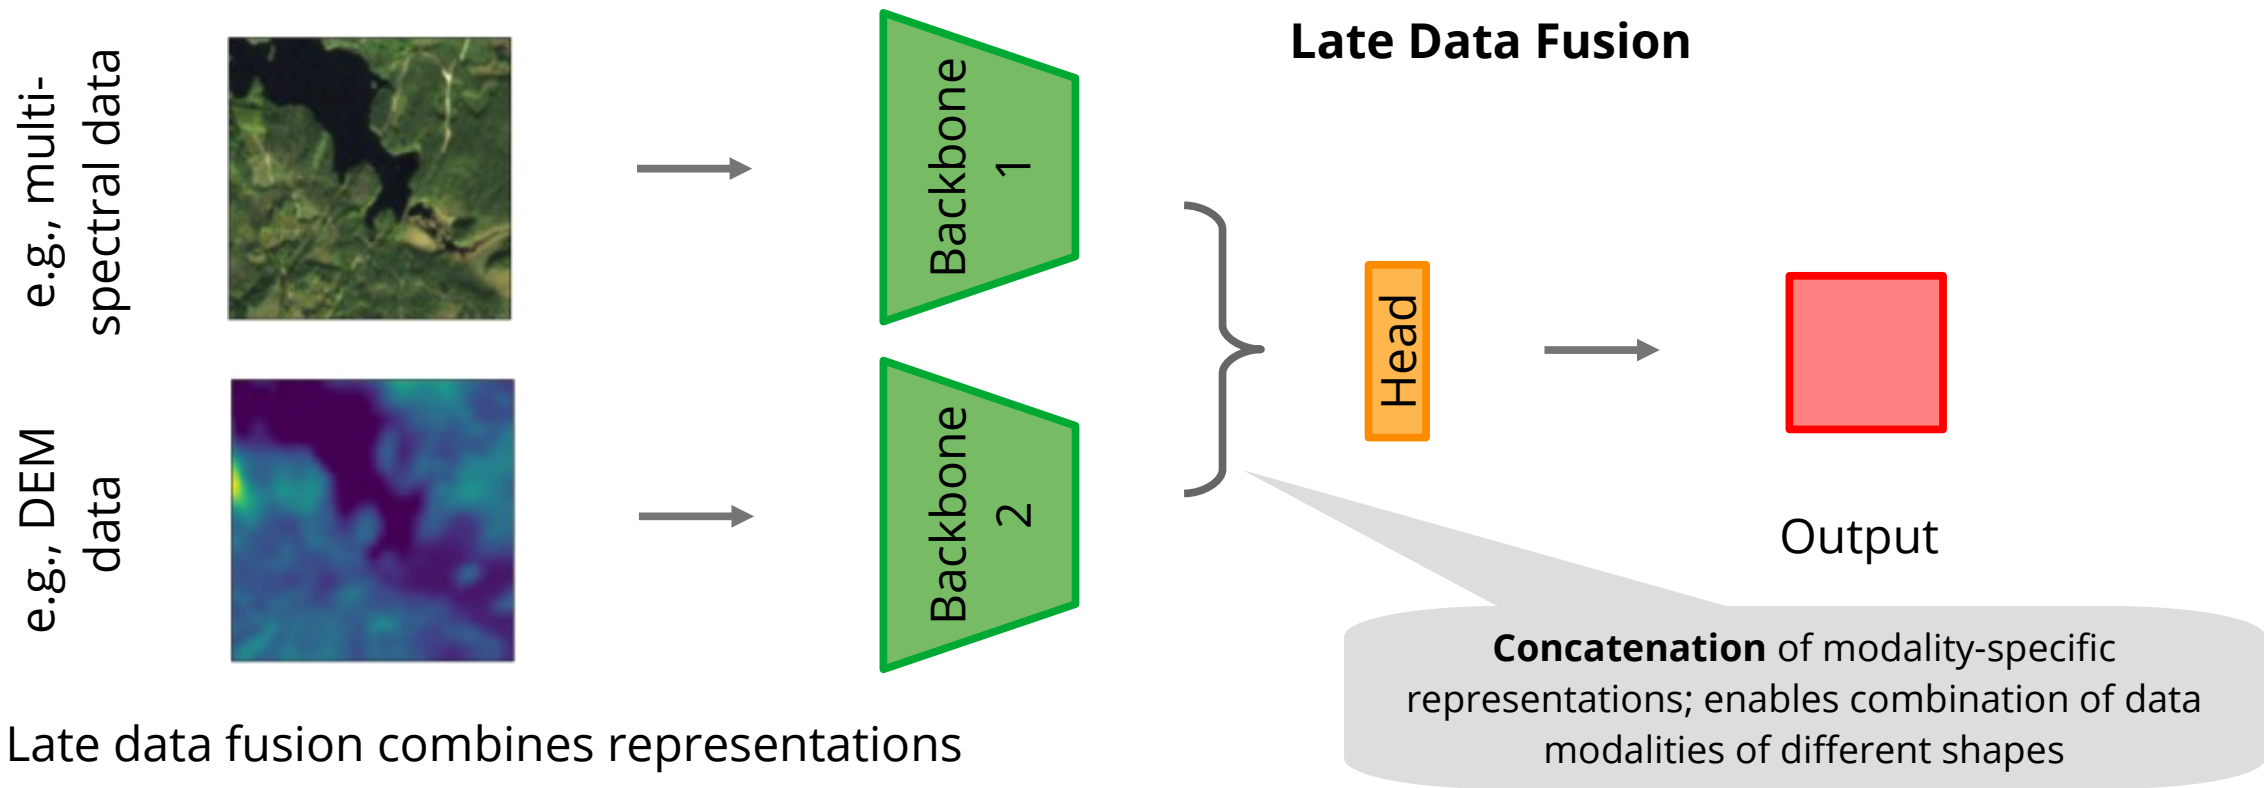


By combining different data modalities early on, we provide additional information to our model.



# How can we use annotated data more efficiently?


We can improve efficiency by leveraging different data modalities in one of two data fusion approaches.



Late data fusion combines representations from different data modalities, representing a more flexible approach for data fusion.

# How can we use annotated data more efficiently?


By forcing the network to learn several tasks at the same time, it will focus on extracting only the most relevant information.



The different heads are trained simultaneously by optimizing the (weighted) sum of the individual losses. As a result, the performance on each task is (typically) better than if trained individually.

# How can we use unannotated data?

We can leverage (large amounts of) unannotated data for pretraining our model on a suitable pretext task.



# Let's get our hands dirty

In the remainder of this tutorial, we will implement and experiment with some of the aforementioned methods and techniques:

- Data Fusion (Michael)
- Multitask Learning (Joëlle)
- Self-supervised Learning (Linus)

We will discuss these techniques in separate Jupyter Notebooks. The overall structure and methods for data handling and other aspects should be consistent between these Notebooks for easier understanding.

Keep in mind that the results provided here are not really representative, as the dataset that we use is tiny. Feel free to use the code from these Notebooks for your own research!